

Adarsh Solanki
As5nr
3/15/12
CS 2150
Post-lab 6

The Big-Theta running complexity of my program can be calculated as follows:

The scanning of the dictionary file and the insertion of the words into the hashTable are a linear function of W (words). The reading and mapping of the grid file is a linear function of R and C (rows and columns). The brute-force search through every combination of grid position ($R * C$) and word length (small constant) can be considered a complexity of $R * C * 8$ (number of directions), which reduces to $R * C$.

This makes the final run-time complexity of the program $W + RC + RC = W + 2RC = \Theta(W + RC)$.

In terms of optimization of performance, I made three improvements: delayed all console outputs, optimized the hash function, and increased the number of buckets.

The most important improvement I made regarded minimizing iterations on the nested for-loops. Functionally my algorithm seemed to stop finding certain edge-cases because of an improper middle condition in one of my for-loops. This led to improper wordCounts upon program termination and missing words because I was erroneously checking that the length was smaller than the rows-currRow, but that

was missing large words going in directions other than East. I decided to store the maximum word size of the hashTable's values as a public data member, and iterated my length from 3 to maxWordSize in order to reduce extra getWordInTable lookups. This ended up increasing the speed of my algorithm by orders of magnitude, reducing my time for 300x300.grid.txt/words2.txt from 52 seconds to 2.36. This was the optimization that was keeping all of my searches so slow all along, as I was doing 1000's of extra calls to hashTable::contains(). This fixed the output for all of the grids to match the given sample output.

On a smaller note, I decided to use a stringstream rather than piping to cout and printed the resulting string only after all of the searching was done and the timer stopped. This increased average speed ~10% for processing a word puzzle. I never knew how computationally expensive console logging could be, but it makes sense, as altering pixels on the display is much more involved than bit manipulation at the memory level.

I improved the hash function by reducing repetition accidentally added due to factoring. I initially had a hash function which made a generic hash and then a myHash function to scale it down to the correct size to fit into the hashTable. I instead merged the two and was able to save a few repeated calculations.

In the constructor of my hashTable, before the primeness of the size argument is checked, I inserted a statement to multiply it by two by using a left bitwise shift. This leads to more overall buckets which reduces collisions and helps make sure each bucket has a smaller average size.

Condition (words.txt // 50x50.grid.txt)	Time (s)
Original application	0.098
Slow hash (hash = str[1] * 78)	3.18
Bad size (no check for prime, not large enough)	0.11

For the slow hash condition, the hash function is terrible because it leads to a LOT of collision. First of all, it only checks one character of the string, so many strings can get similar results. Also, it doesn't even multiply by a prime constant, so there are multiple factors that can lead to the same hashCode product, leading to yet more collisions.

For the bad size, I simply removed my bitshift and also the check for primeness. Without a prime number of buckets, the modulus function is repeated often, for example all of the factors of a number evaluate to 0. This increases collisions. The removal of the bitshift just leads to way fewer buckets, which increases the need for separate chaining and the load factor.

In the end, my output for words2.txt/300x300.grid.txt ended up being 2.36399 seconds, significant increase from the 52 seconds it took my code to complete prior to my optimizations. This is a speed increase by a factor of $52/2.36 \sim 22x$! This is an astronomical algorithmic simplification.